

Applying Operating System Principles to SDN Controller Design

Oliver Michel, Matthew Monaco, Eric Keller

Invited Talk University of Illinois at Urbana-Champaign
April 11th, 2014.



University
of Colorado
Boulder



“What we clearly need is an “operating system” for networks, one that provides a uniform and centralized programmatic interface to the entire network.”

[Gude et.al.’08]

NOX: Towards an Operating System for Networks

Natasha Gude
Nicira Networks

Ben Pfaff
Nicira Networks

Teemu Koponen
HIIT

Martín Casado
Nicira Networks

Scott Shenker
University of California,
Berkeley

Justin Pettit
Nicira Networks

Nick McKeown
Stanford University

This article is an editorial note submitted to CCR. It has NOT been peer reviewed.
Authors take full responsibility for this article’s technical content.
Comments can be posted through CCR Online.

Categories and Subject Descriptors

C.2.1 [Computer-Communication Networks]: Network Architecture and Design

General Terms

Design, Experimentation, Performance

Keywords

Architecture, Management, Network, Security

1 Introduction

As anyone who has operated a large network can attest, enterprise networks are difficult to manage. That they have

interface to the entire network.¹ Analogous to the read and write access to various resources provided by computer operating systems, a network operating system provides the ability to *observe* and *control* a network.

A network operating system does not manage the network itself; it merely provides a programmatic interface. *Applications* implemented on top of the network operating system perform the actual management tasks.² The programmatic interface should be general enough to support a broad spectrum of network management applications.

Such a network operating system represents two major conceptual departures from the status quo. First, the network operating system presents programs with a *centralized* programming model³; programs are written as if the entire network were present on a single machine (*i.e.*, one would use Dijkstra to compute shortest paths, not Bellman-Ford). This



Extend an existing operating system
and its user space software ecosystem
in order to serve as a practical
network operating system



Distributions



Software Projects





\$



```
$
```

```
$ grep  
$ find  
$ sed  
$ sort  
$ cat
```



```
$ cd /var/log  
$ grep that_annoying_problem syslog
```



```
$ cd /var/log  
$ grep that_annoying_problem syslog
```


```
$ ps -A | sort -k3nr | head -10  
$ kill -TERM 12345
```



```
#!/bin/bash

procs=(  
    ps -A | sort -k3nr | head -$1 | tr -s ' ' \  
    | cut -d' ' -f9  
)  
  
for p in "${procs[@]}"; do  
    printf "%d is misbehaving\n" "$p" >&2 kill -TERM "$p"  
    sleep 3  
    kill -KILL "$p"  
done
```





















C++





C++



Python





C++



Python



Python





C++



Python



Python



Java





C++



Python



Python



Java



Java





C++



Python



Python




Java



Java



Ruby/C



yanc

yet another network controller



yanc


yet another network controller

```
/net
  hosts
  switches
    sw1
    sw2
  views
    http
    management-net
      hosts
      switches
      views
```




yanc

yet another network controller



yanc


yet another network controller



Why a filesystem?



Logically Distinct Applications




Independent Development

```
# apt-get install yanc-learning-switch  
# apt-get install yanc-router
```

```
$ git clone git@github.com:yanc/yanc-fw  
$ cd yanc-fw  
$ make  
# make install
```



Any Programming Language




Yanc Components

- ▶ The Yanc Filesystem
- ▶ Hardware Decoupling
- ▶ Operating System Integration
- ▶ Dynamic Application Composition
- ▶ Distribution



The Yanc Filesystem



The Yanc Filesystem

```
01:02:03:04:05:06/
  └── flows/
  └── packet_in/
  └── packet_out/
  └── ports/
  └── actions
  └── capabilities
  └── datapath_id
  └── flags
  └── miss_send_len
  └── n_buffers
  └── n_tables
```




The Yanc Filesystem

```
port_01/
  config.flood
  config.fwd
  config.packet_in
  config.port_down
  config.recv
  config.recv_stp
  config.stp
  hw_addr
  peer -> /dev/null
  port_no
  state.link_up
  state.stp_forward
  state.stp_learn
  state.stp_listen
  stats.collisions
  stats.rx_bytes
  stats.rx_crc_err
  stats.rx_dropped
  stats.rx_errors
  stats.rx_frame_err
  stats.rx_over_err
  stats.rx_packets
  stats.tx_bytes
  stats.tx_dropped
  stats.tx_errors
  stats.tx_packets
```



The Yanc Filesystem



The Yanc Filesystem

```
arp_flow
  └── counters/
      ├── match.dl_type
      ├── match.dl_src
      ├── action.out
      ├── priority
      ├── timeout
      └── version
```



The Yanc Filesystem

```
$ echo 1 > port_1.port_down
```

```
$ mkdir -p switches/00:01:02:0a:0b:0c/flows/my_flow
```



The Yanc Filesystem

```
$ find /net -name tp_dst -exec grep 22 {} +
```

```
#!/bin/bash
flowdir=/net/switches/"$1"/flows/"$2"
mkdir "$flowdir"
echo ff:ff:ff:ff:ff:ff > "$flowdir"/match.dl_dst
echo 0x0806 > "$flowdir"/match.dl_type
echo FL00D > "$flowdir"/action.out
```



Hardware Decoupling and Middleboxes

- ▶ Support for different physical devices and device classes
 - Device Drivers
 - Schema-based filesystem




Device Drivers




Device Drivers

Applications




Device Drivers



Device Drivers

Applications



Filesystem

/net/switch1/flow1/...

/net/firewall1/rule1/...


/net/firewall2/rule1/...

/net/ids/alerts/...

Drivers




Device Drivers



Schema-based Filesystem

```
nodes {  
    switch-* {  
        u64 id;  
        ports/* {  
            ln peer;  
        }  
    }  
  
    host-* {  
        eth mac_addr;  
        str hostname;  
        ln port-*;  
    }  
}
```



Operating System Integration

- ▶ Integrate tightly with Linux
- ▶ Use off-the-shelf technologies
- ▶ Encourage active ecosystem




Operating System Integration

What can we use from a traditional operating system?



Operating System Integration


What can we use from a traditional operating system?



Operating System Integration

What can we use from a traditional operating system?

- ▶ Inotify
- ▶ File Permissions and ACLs
- ▶ Namespaces and Control Groups
- ▶ Layered Filesystem



Application Composition

- ▶ Today's controllers show mostly monolithic design
- ▶ Distinct applications must be composed



Composition in Yanc

- ▶ Frenetic/Pyretic introduced Parallel/Sequential Composition
 - per packet abstraction
 - compiled into single application
- ▶ Yanc controls read/writes to the network state
- ▶ Yanc provides dynamic composition



Motivating Example



Motivating Example

Fast Router



Motivating Example

Fast Router

Route Optimizer



Motivating Example

Fast Router

Route Optimizer

Link Tapper



Motivating Example

Fast Router

Route Optimizer

Link Tapper



Motivating Example

Fast Router

Route Optimizer

Link Tapper



Motivating Example




Motivating Example




Motivating Example



Optimizing a flow



Adding a Linktapper




Distribution


- ▶ Filesystem representation allows layering
 - Distributed filesystem on top of yancfs
- ▶ Yanc may run on multiple servers or even switches




Layered Filesystem




Locking




Locking




Locking




Locking




Locking




Locking




Locking




Locking




Locking




Locking




Locking



Locking



Locking



Prototype

- ▶ Filesystem implemented using FUSE (C)
- ▶ OpenFlow Driver (C++)
- ▶ Applications (Python and Ruby)
 - Topology Discovery
 - Learning Switch
 - Fast Router
 - Route Optimizer
 - Link Tapper
- ▶ several Bash management scripts and static flow pusher



Conclusion

- ▶ Defined a Universal Interface for Network Control
- ▶ Allow Control Applications as Separate Processes
- ▶ Implemented Interface
- ▶ Leveraging Existing OS Technologies
- ▶ Built Functional Applications on top of Interface



Thank you!

Matthew Monaco, Oliver Michel, and Eric Keller. Applying operating system principles to SDN controller design. In *Proceedings of the Twelfth ACM Workshop on Hot Topics in Networks (HotNets-XII)*. 2013.



Backup Slides



Programmability

```
#!/usr/bin/env python3

def new_switch(id, n_tables=1):
    pass

def write_flow(switch, matches=[], actions=[]):
    pass
```


```
#ifndef _YANC_H_
#define _YANC_H_

int new_switch(uint64_t, uint8_t);
int write_flow(const char* path, match_t*, action_t*);


#endif/*_YANC_H_*/
```



Architecture



Performance



file system	op	ops/s
yancfs	mkdir	10245
yancfs	rmdir	11192
tmpfs	mkdir	721877
tmpfs	rmdir	782249

Shared Library

